



Sixth Edition

# Database System Concepts

Abraham Silberschatz • Henry F. Korth • S. Sudarshan

# DATABASE SYSTEM CONCEPTS

SIXTH EDITION

Abraham Silberschatz

*Yale University*

Henry F. Korth

*Lehigh University*

S. Sudarshan

*Indian Institute of Technology, Bombay*





## DATABASE SYSTEM CONCEPTS, SIXTH EDITION

Published by McGraw-Hill, a business unit of The McGraw-Hill Companies, Inc., 1221 Avenue of the Americas, New York, NY 10020. Copyright © 2011 by The McGraw-Hill Companies, Inc. All rights reserved. Previous editions © 2006, 2002, and 1999. No part of this publication may be reproduced or distributed in any form or by any means, or stored in a database or retrieval system, without the prior written consent of The McGraw-Hill Companies, Inc., including, but not limited to, in any network or other electronic storage or transmission, or broadcast for distance learning.

Some ancillaries, including electronic and print components, may not be available to customers outside the United States.

This book is printed on acid-free paper.

1 2 3 4 5 6 7 8 9 0 DOC/DOC 1 0 9 8 7 6 5 4 3 2 1 0

ISBN 978-0-07-352332-3

MHID 0-07-352332-1

Global Publisher: *Raghothaman Srinivasan*

Director of Development: *Kristine Tibbetts*

Senior Marketing Manager: *Curt Reynolds*

Project Manager: *Melissa M. Leick*

Senior Production Supervisor: *Laura Fuller*

Design Coordinator: *Brenda A. Rolwes*

Cover Designer: *Studio Montage, St. Louis, Missouri*

(USE) Cover Image: © *Brand X Pictures/PunchStock*

Compositor: *Aptara®, Inc.*

Typeface: *10/12 Palatino*

Printer: *R. R. Donnelley*

All credits appearing on page or at the end of the book are considered to be an extension of the copyright page.

### Library of Congress Cataloging-in-Publication Data

Silberschatz, Abraham.

Database system concepts / Abraham Silberschatz. — 6th ed.

p. cm.

ISBN 978-0-07-352332-3 (alk. paper)

1. Database management. I. Title.

QA76.9.D3S5637 2011

005.74—dc22

2009039039

The Internet addresses listed in the text were accurate at the time of publication. The inclusion of a Web site does not indicate an endorsement by the authors of McGraw-Hill, and McGraw-Hill does not guarantee the accuracy of the information presented at these sites.

*In memory of my father Joseph Silberschatz  
my mother Vera Silberschatz  
and my grandparents Stepha and Aaron Rosenblum*

*Avi Silberschatz*

*To my wife, Joan  
my children, Abigail and Joseph  
and my parents, Henry and Frances*

*Hank Korth*

*To my wife, Sita  
my children, Madhur and Advaith  
and my mother, Indira*

*S. Sudarshan*

*This page intentionally left blank*

---

# Contents

## Chapter 1 Introduction

- 1.1 Database-System Applications 1
- 1.2 Purpose of Database Systems 3
- 1.3 View of Data 6
- 1.4 Database Languages 9
- 1.5 Relational Databases 12
- 1.6 Database Design 15
- 1.7 Data Storage and Querying 20
- 1.8 Transaction Management 22
- 1.9 Database Architecture 23
- 1.10 Data Mining and Information Retrieval 25
- 1.11 Specialty Databases 26
- 1.12 Database Users and Administrators 27
- 1.13 History of Database Systems 29
- 1.14 Summary 31
- Exercises 33
- Bibliographical Notes 35

## PART ONE ■ RELATIONAL DATABASES

## Chapter 2 Introduction to the Relational Model

- 2.1 Structure of Relational Databases 39
- 2.2 Database Schema 42
- 2.3 Keys 45
- 2.4 Schema Diagrams 46
- 2.5 Relational Query Languages 47
- 2.6 Relational Operations 48
- 2.7 Summary 52
- Exercises 53
- Bibliographical Notes 55

## Chapter 3 Introduction to SQL

- 3.1 Overview of the SQL Query Language 57
- 3.2 SQL Data Definition 58
- 3.3 Basic Structure of SQL Queries 63
- 3.4 Additional Basic Operations 74
- 3.5 Set Operations 79
- 3.6 Null Values 83
- 3.7 Aggregate Functions 84
- 3.8 Nested Subqueries 90
- 3.9 Modification of the Database 98
- 3.10 Summary 104
- Exercises 105
- Bibliographical Notes 112

**Chapter 4 Intermediate SQL**

- 4.1 Join Expressions 113
- 4.2 Views 120
- 4.3 Transactions 127
- 4.4 Integrity Constraints 128
- 4.5 SQL Data Types and Schemas 136
- 4.6 Authorization 143
- 4.7 Summary 150
  - Exercises 152
  - Bibliographical Notes 156

**Chapter 5 Advanced SQL**

- 5.1 Accessing SQL From a Programming Language 157
- 5.2 Functions and Procedures 173
- 5.3 Triggers 180
- 5.4 Recursive Queries\*\* 187
- 5.5 Advanced Aggregation Features\*\* 192
- 5.6 OLAP\*\* 197
- 5.7 Summary 209
  - Exercises 211
  - Bibliographical Notes 216

**Chapter 6 Formal Relational Query Languages**

- 6.1 The Relational Algebra 217
- 6.2 The Tuple Relational Calculus 239
- 6.3 The Domain Relational Calculus 245
- 6.4 Summary 248
  - Exercises 249
  - Bibliographical Notes 254

**PART TWO ■ DATABASE DESIGN****Chapter 7 Database Design and the E-R Model**

- 7.1 Overview of the Design Process 259
- 7.2 The Entity-Relationship Model 262
- 7.3 Constraints 269
- 7.4 Removing Redundant Attributes in Entity Sets 272
- 7.5 Entity-Relationship Diagrams 274
- 7.6 Reduction to Relational Schemas 283
- 7.7 Entity-Relationship Design Issues 290
- 7.8 Extended E-R Features 295
- 7.9 Alternative Notations for Modeling Data 304
- 7.10 Other Aspects of Database Design 310
- 7.11 Summary 313
  - Exercises 315
  - Bibliographical Notes 321

## Chapter 8 Relational Database Design

- 8.1 Features of Good Relational Designs 323
- 8.2 Atomic Domains and First Normal Form 327
- 8.3 Decomposition Using Functional Dependencies 329
- 8.4 Functional-Dependency Theory 338
- 8.5 Algorithms for Decomposition 348
- 8.6 Decomposition Using Multivalued Dependencies 355
- 8.7 More Normal Forms 360
- 8.8 Database-Design Process 361
- 8.9 Modeling Temporal Data 364
- 8.10 Summary 367
- Exercises 368
- Bibliographical Notes 374

## Chapter 9 Application Design and Development

- 9.1 Application Programs and User Interfaces 375
- 9.2 Web Fundamentals 377
- 9.3 Servlets and JSP 383
- 9.4 Application Architectures 391
- 9.5 Rapid Application Development 396
- 9.6 Application Performance 400
- 9.7 Application Security 402
- 9.8 Encryption and Its Applications 411
- 9.9 Summary 417
- Exercises 419
- Bibliographical Notes 426

# PART THREE ■ DATA STORAGE AND QUERYING

## Chapter 10 Storage and File Structure

- 10.1 Overview of Physical Storage Media 429
- 10.2 Magnetic Disk and Flash Storage 432
- 10.3 RAID 441
- 10.4 Tertiary Storage 449
- 10.5 File Organization 451
- 10.6 Organization of Records in Files 457
- 10.7 Data-Dictionary Storage 462
- 10.8 Database Buffer 464
- 10.9 Summary 468
- Exercises 470
- Bibliographical Notes 473

## Chapter 11 Indexing and Hashing

- 11.1 Basic Concepts 475
- 11.2 Ordered Indices 476
- 11.3 B<sup>+</sup>-Tree Index Files 485
- 11.4 B<sup>+</sup>-Tree Extensions 500
- 11.5 Multiple-Key Access 506
- 11.6 Static Hashing 509
- 11.7 Dynamic Hashing 515
- 11.8 Comparison of Ordered Indexing and Hashing 523
- 11.9 Bitmap Indices 524
- 11.10 Index Definition in SQL 528
- 11.11 Summary 529
- Exercises 532
- Bibliographical Notes 536



**Chapter 12 Query Processing**

- 12.1 Overview 537
- 12.2 Measures of Query Cost 540
- 12.3 Selection Operation 541
- 12.4 Sorting 546
- 12.5 Join Operation 549
- 12.6 Other Operations 563
- 12.7 Evaluation of Expressions 567
- 12.8 Summary 572
- Exercises 574
- Bibliographical Notes 577

**Chapter 13 Query Optimization**

- 13.1 Overview 579
- 13.2 Transformation of Relational Expressions 582
- 13.3 Estimating Statistics of Expression Results 590
- 13.4 Choice of Evaluation Plans 598
- 13.5 Materialized Views\*\* 607
- 13.6 Advanced Topics in Query Optimization\*\* 612
- 13.7 Summary 615
- Exercises 617
- Bibliographical Notes 622

**PART FOUR ■ TRANSACTION MANAGEMENT****Chapter 14 Transactions**

- 14.1 Transaction Concept 627
- 14.2 A Simple Transaction Model 629
- 14.3 Storage Structure 632
- 14.4 Transaction Atomicity and Durability 633
- 14.5 Transaction Isolation 635
- 14.6 Serializability 641
- 14.7 Transaction Isolation and Atomicity 646
- 14.8 Transaction Isolation Levels 648
- 14.9 Implementation of Isolation Levels 650
- 14.10 Transactions as SQL Statements 653
- 14.11 Summary 655
- Exercises 657
- Bibliographical Notes 660

**Chapter 15 Concurrency Control**

- 15.1 Lock-Based Protocols 661
- 15.2 Deadlock Handling 674
- 15.3 Multiple Granularity 679
- 15.4 Timestamp-Based Protocols 682
- 15.5 Validation-Based Protocols 686
- 15.6 Multiversion Schemes 689
- 15.7 Snapshot Isolation 692
- 15.8 Insert Operations, Delete Operations, and Predicate Reads 697
- 15.9 Weak Levels of Consistency in Practice 701
- 15.10 Concurrency in Index Structures\*\* 704
- 15.11 Summary 708
- Exercises 712
- Bibliographical Notes 718

## Chapter 16 Recovery System

- 16.1 Failure Classification 721
- 16.2 Storage 722
- 16.3 Recovery and Atomicity 726
- 16.4 Recovery Algorithm 735
- 16.5 Buffer Management 738
- 16.6 Failure with Loss of Nonvolatile Storage 743
- 16.7 Early Lock Release and Logical Undo Operations 744
- 16.8 ARIES\*\* 750
- 16.9 Remote Backup Systems 756
- 16.10 Summary 759
  - Exercises 762
  - Bibliographical Notes 766

## PART FIVE ■ SYSTEM ARCHITECTURE

### Chapter 17 Database-System Architectures

- 17.1 Centralized and Client–Server Architectures 769
- 17.2 Server System Architectures 772
- 17.3 Parallel Systems 777
- 17.4 Distributed Systems 784
- 17.5 Network Types 788
- 17.6 Summary 791
  - Exercises 793
  - Bibliographical Notes 794

### Chapter 18 Parallel Databases

- 18.1 Introduction 797
- 18.2 I/O Parallelism 798
- 18.3 Interquery Parallelism 802
- 18.4 Intraquery Parallelism 803
- 18.5 Intraoperation Parallelism 804
- 18.6 Interoperation Parallelism 813
- 18.7 Query Optimization 814
- 18.8 Design of Parallel Systems 815
- 18.9 Parallelism on Multicore Processors 817
- 18.10 Summary 819
  - Exercises 821
  - Bibliographical Notes 824

### Chapter 19 Distributed Databases

- 19.1 Homogeneous and Heterogeneous Databases 825
- 19.2 Distributed Data Storage 826
- 19.3 Distributed Transactions 830
- 19.4 Commit Protocols 832
- 19.5 Concurrency Control in Distributed Databases 839
- 19.6 Availability 847
- 19.7 Distributed Query Processing 854
- 19.8 Heterogeneous Distributed Databases 857
- 19.9 Cloud-Based Databases 861
- 19.10 Directory Systems 870
- 19.11 Summary 875
  - Exercises 879
  - Bibliographical Notes 883

## PART SIX ■ DATA WAREHOUSING, DATA MINING, AND INFORMATION RETRIEVAL

### Chapter 20 Data Warehousing and Mining

- |                                      |                                     |
|--------------------------------------|-------------------------------------|
| 20.1 Decision-Support Systems 887    | 20.7 Clustering 907                 |
| 20.2 Data Warehousing 889            | 20.8 Other Forms of Data Mining 908 |
| 20.3 Data Mining 893                 | 20.9 Summary 909                    |
| 20.4 Classification 894              | Exercises 911                       |
| 20.5 Association Rules 904           | Bibliographical Notes 914           |
| 20.6 Other Types of Associations 906 |                                     |

### Chapter 21 Information Retrieval

- |   |   |
|---|---|
| 21.1 Overview 915                           | 21.7 Crawling and Indexing the Web 930                  |
| 21.2 Relevance Ranking Using Terms 917      | 21.8 Information Retrieval: Beyond Ranking of Pages 931 |
| 21.3 Relevance Using Hyperlinks 920         | 21.9 Directories and Categories 935                     |
| 21.4 Synonyms, Homonyms, and Ontologies 925 | 21.10 Summary 937                                       |
| 21.5 Indexing of Documents 927              | Exercises 939   |
| 21.6 Measuring Retrieval Effectiveness 929  | Bibliographical Notes 941                               |

## PART SEVEN ■ SPECIALTY DATABASES

### Chapter 22 Object-Based Databases

- |   |  |
|---|--|
| 22.1 Overview 945                                   | 22.8 Persistent Programming Languages 964          |
| 22.2 Complex Data Types 946                         | 22.9 Object-Relational Mapping 973                 |
| 22.3 Structured Types and Inheritance in SQL 949    | 22.10 Object-Oriented versus Object-Relational 973 |
| 22.4 Table Inheritance 954                          | 22.11 Summary 975                                  |
| 22.5 Array and Multiset Types in SQL 956            | Exercises 976                                      |
| 22.6 Object-Identity and Reference Types in SQL 961 | Bibliographical Notes 980                          |
| 22.7 Implementing O-R Features 963                  |  |

### Chapter 23 XML

- |   |                               |
|---|-------------------------------|
| 23.1 Motivation 981                             | 23.6 Storage of XML Data 1009 |
| 23.2 Structure of XML Data 986                  | 23.7 XML Applications 1016    |
| 23.3 XML Document Schema 990                    | 23.8 Summary 1019             |
| 23.4 Querying and Transformation 998            | Exercises 1021                |
| 23.5 Application Program Interfaces to XML 1008 | Bibliographical Notes 1024    |

## PART EIGHT ■ ADVANCED TOPICS

### Chapter 24 Advanced Application Development

- |  |      |                       |      |
|--|------|-----------------------|------|
| 24.1 Performance Tuning                      | 1029 | 24.4 Standardization  | 1051 |
| 24.2 Performance Benchmarks                  | 1045 | 24.5 Summary          | 1056 |
| 24.3 Other Issues in Application Development | 1048 | Exercises             | 1057 |
|  |      | Bibliographical Notes | 1059 |

### Chapter 25 Spatial and Temporal Data and Mobility

- |                                  |      |                                      |      |
|----------------------------------|------|--------------------------------------|------|
| 25.1 Motivation                  | 1061 | 25.5 Mobility and Personal Databases | 1079 |
| 25.2 Time in Databases           | 1062 | 25.6 Summary                         | 1085 |
| 25.3 Spatial and Geographic Data | 1064 | Exercises                            | 1087 |
| 25.4 Multimedia Databases        | 1076 | Bibliographical Notes                | 1089 |

### Chapter 26 Advanced Transaction Processing

- |                                      |      |                                 |      |
|--------------------------------------|------|---------------------------------|------|
| 26.1 Transaction-Processing Monitors | 1091 | 26.6 Long-Duration Transactions | 1109 |
| 26.2 Transactional Workflows         | 1096 | 26.7 Summary                    | 1115 |
| 26.3 E-Commerce                      | 1102 | Exercises                       | 1117 |
| 26.4 Main-Memory Databases           | 1105 | Bibliographical Notes           | 1119 |
| 26.5 Real-Time Transaction Systems   | 1108 |                                 |      |

## PART NINE ■ CASE STUDIES

### Chapter 27 PostgreSQL

- |   |      |  |      |
|---|------|--|------|
| 27.1 Introduction                         | 1123 | 27.5 Storage and Indexing              | 1146 |
| 27.2 User Interfaces                      | 1124 | 27.6 Query Processing and Optimization | 1151 |
| 27.3 SQL Variations and Extensions        | 1126 | 27.7 System Architecture               | 1154 |
| 27.4 Transaction Management in PostgreSQL | 1137 | Bibliographical Notes                  | 1155 |

### Chapter 28 Oracle

- |   |      |   |      |
|---|------|---|------|
| 28.1 Database Design and Querying Tools | 1157 | 28.6 System Architecture                          | 1183 |
| 28.2 SQL Variations and Extensions      | 1158 | 28.7 Replication, Distribution, and External Data | 1188 |
| 28.3 Storage and Indexing               | 1162 | 28.8 Database Administration Tools                | 1189 |
| 28.4 Query Processing and Optimization  | 1172 | 28.9 Data Mining                                  | 1191 |
| 28.5 Concurrency Control and Recovery   | 1180 | Bibliographical Notes                             | 1191 |

**Chapter 29   IBM DB2 Universal Database**

29.1 Overview	1193	29.9 Tools and Utilities	1215
29.2 Database-Design Tools	1194	29.10 Concurrency Control and Recovery	1217
29.3 SQL Variations and Extensions	1195	29.11 System Architecture	1219
29.4 Storage and Indexing	1200	29.12 Replication, Distribution, and External Data	1220
29.5 Multidimensional Clustering	1203	29.13 Business Intelligence Features	1221
29.6 Query Processing and Optimization	1207	Bibliographical Notes	1222
29.7 Materialized Query Tables	1212		
29.8 Autonomic Features in DB2	1214		

**Chapter 30   Microsoft SQL Server**

30.1 Management, Design, and Querying Tools	1223	30.8 Distributed Heterogeneous Query Processing	1250
30.2 SQL Variations and Extensions	1228	30.9 Replication	1251
30.3 Storage and Indexing	1233	30.10 Server Programming in .NET	1253
30.4 Query Processing and Optimization	1236	30.11 XML Support	1258
30.5 Concurrency and Recovery	1241	30.12 SQL Server Service Broker	1261
30.6 System Architecture	1246	30.13 Business Intelligence	1263
30.7 Data Access	1248	Bibliographical Notes	1267

**PART TEN   ■   APPENDICES**

**Appendix A   Detailed University Schema**

A.1 Full Schema	1271	A.3 Sample Data	1276
A.2 DDL	1272		

**Appendix B   Advanced Relational Design (contents online)**

B.1 Multivalued Dependencies	B1	Exercises	B10
B.3 Domain-Key Normal Form	B8	Bibliographical Notes	B12
B.4 Summary	B10		

**Appendix C   Other Relational Query Languages (contents online)**

C.1 Query-by-Example	C1	C.4 Summary	C25
C.2 Microsoft Access	C9	Exercises	C26
C.3 Datalog	C11	Bibliographical Notes	C30

## Appendix D Network Model (contents online)

D.1 Basic Concepts	D1	D.6 DBTG Set-Processing Facility	D22
D.2 Data-Structure Diagrams	D2	D.7 Mapping of Networks to Files	D27
D.3 The DBTG CODASYL Model	D7	D.8 Summary	D31
D.4 DBTG Data-Retrieval Facility	D13	Exercises	D32
D.5 DBTG Update Facility	D20	Bibliographical Notes	D35

## Appendix E Hierarchical Model (contents online)

E.1 Basic Concepts	E1	E.6 Mapping of Hierarchies to Files	E22
E.2 Tree-Structure Diagrams	E2	E.7 The IMS Database System	E24
E.3 Data-Retrieval Facility	E13	E.8 Summary	E25
E.4 Update Facility	E17	Exercises	E26
E.5 Virtual Records	E20	Bibliographical Notes	E29

## Bibliography 1283

## Index 1315

*This page intentionally left blank*